

Waldemar Quevedo

INTERESTS

PaaS & Distributed Systems, Web Operations, Literate programming

WORK EXPERIENCE

JAN 2012 - *Present*

PaaS Developer and Administrator, [Rakuten](#)

Joined the Rakuten PaaS project as one of the first members of the team. The project started as an initiative to improve the development cycle of services by automating away the provisioning of resources and providing a deployment pipeline for application developers, basing on the Cloudfoundry project from VMware.

RESPONSIBILITIES

- Development of features for making the PaaS compliant with what Rakuten users needed
- Maintenance and operation of the PaaS
 - Ruby runtime & Rails support
 - Infrastructure as code approach with Chef
 - Automated provisioning of resources with vSphere as the virtualization environment
 - Capacity management and coordination with infrastructure teams
 - Setup of Cloudfoundry deployments in new datacenters
- 24/7 on call rotation to support issues in the system and users of the platform
- Evaluation of next generation PaaS platforms

LOGGING RELATED PROJECTS

Designed and implemented the log collection infrastructure to be used throughout the PaaS by using Fluentd, given that there was no consolidated logging approach at the time in Cloudfoundry.

Application logs collection and streaming: Modified the Cloudfoundry DEA component so that it would attach a background Fluentd process dynamically to an application during boot, so that it would persist any generated logs by forwarding to a log aggregation layer, as well as enabling realtime 'tailing' support of the log stream through the CLI.

Developed plugins to tag the logs emitted by the users application in realtime and the Ruby logging library which would connect transparently to the log collector port from within the application.

Log forwarding endpoints integration: Collaborated with Hadoop and Splunk teams to re-route the collected logs from the PaaS to send to their environments.

Defined the SLAs needed to clarify the responsibilities from these collaborations (log delay, buffering size) as well as writing Fluentd plugins to normalize the emitted logs so that it would fit the format expected by those teams.

This feature ended up being one of the selling points of the system since it saved application teams having to manage their own collectors in order to send to Hadoop.

Logs routing as a service: Enhanced the architecture so that log forwarding bindings could be created dynamically by users. This way a user could define through an API

whether the access and application logs generated by an application should be routed to other endpoints besides the default persistence storage.

The mechanism here was to have a Fluentd polling the configuration about the logging service bindings from a user (approach done by Cloudfoundry services components) to create forwarding rules to route the emit to the defined endpoint.

OCT 2010 - DEC
2011

Application Engineer, [Rakuten](#)

Worked as part of the [Rakuten Research](#) Development Team. This is a service which allows registered Rakuten members to take part in marketing research surveys in exchange of Rakuten points that can be used on the main Rakuten Ichiba shopping site and other parts of the Rakuten ecosystem.

The main tasks involved were the maintenance and operation of the Rakuten Research systems, as well as adding new features to the product.

This time was very early on in Rakuten efforts of globalization, so all communication and development was handled in Japanese.

APR - AUG, 2010

Google Summer of Code Student, [OpenStreetMap](#)

Through the Google Summer of Code program, I worked on a web application for OpenStreetMap to create and customize printable maps using spatial data. The main technologies used were Django, Twisted and Mapnik.

MAY - DEC, 2009

Backend Developer, [Inovaz](#)

Worked as an intern at Inovaz, a software startup partly based in Guadalajara, Mexico and London (nowadays they go under the name of [Wepow](#)). I took part of the backend team working with Rails 2.3. The project I was working on was [FatStudent](#), a platform where university students from London could apply for part-time jobs using a video CV. The application went live during my internship and turned into a successful launch for students from the University College London.

EDUCATION

OCT - DEC 2014

Hacker School in NYC

Joined the Fall 2014 batch of [HackerSchool](#), a 3 month retreat for programmers based in New York.

I mostly focused on learning about parsers to create new implementations of Org mode which can be supported in environments outside of Emacs, and practical literate programming using Org Babel.

2004 - 2010

INSTITUTION

Bachelor in Computer Science and Technologies

Monterrey Institute of Technology and Higher Education, campus Guadalajara, Mexico

APR - AUG, 2007

INSTITUTION

Exchange Student

Ritsumeikan University. Kyoto, Japan

ACTIVITIES

ORG MODE RUBY
PARSER
MAINTAINER

Maintainer of the Org mode parser used in Github for rendering `.org` files into HTML ([link](#)). Also developed other proof of concept tools that base on the literate programming functionalities from Org mode to be able to run code blocks without using Emacs. ([Screencast](#)).

TECHNICAL SKILLS

PROGRAMMING
LANGUAGES Most comfortable with Ruby, Python, JS & Golang.

PAAS Cloudfoundry. Familiar with the ecosystem around Mesos.

GIS OpenStreetMap Toolchain: PostGIS, Mapnik

LANGUAGES

SPANISH Native

ENGLISH Proficient (TOEFL 617 / TOEIC 880)

JAPANESE Proficient (Business Level, JLPT 2)

CHINESE, FRENCH Beginner

PERSONAL INFORMATION

DATE OF BIRTH May 4th, 1986

NATIONALITY Mexican

PHONE +81 80 4098 0884

E-MAIL waldemar.quevedo@gmail.com

SKYPE waldemarquevedo

BLOG

- <http://literateinfrastructure.org>

CODE

- <https://github.com/wallyqs>